

▼ REAL ESTATE HOUSE PRICE PREDICTION

Objective

- Predict the price of houses from the given dataset by applying different regression based machine learning algorithms.
- Applying different performance boosting methods like feature selection, Hyper parameter tuning etc.
- Utilizing different EDA tools for visualization and data manipulation
- Comparing the performance of the different regression models

About the Dataset

1. title
- Shows the number of bedrooms if it is house. and the details of the location
2. price
- price of the house or the plot
3. size
- size of the house in square feet
4. price_per_sqft
- per square feet price for the house or plot
5. status
- whether the place is open for living or under construction

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

df=pd.read_excel('/content/raw_data.xlsx')
df.head()
```

	Unnamed: 0	title	price	size	price_per_sqft	status
0	0	3 BHK Apartment in CasaGrand Casagrاند Meridian	1.18 Cr	2037	5,793 / sq ft	Under Construction
1	1	2 BHK Apartment in Shree Nandana Elite	57.56 L	1240	4,642 / sq ft	Under Construction
2	2	Residential Plot in Saroj Whispering Winds	32.3 L	2000	1,615 / sq ft	New
3	3	2 BHK Apartment in Shree Nandana Elite	1.33 Cr	1240	10,726 / sq ft	Under Construction

```
df.sample(5)
```

	Unnamed: 0	title	price	size	price_per_sqft	status
56235	56234	\n Salarpuria Gold Summit\n3 BHK Flat\nHennur\n	165 Cr	1933 Sq.Ft.	NaN	NaN
41597	41597	2 BHK Apartment in Abhee nandana	61.9 L	1140	5,430 / sq ft	Ready to move
7161	7161	2 BHK Apartment in G Corp The Icon	1.2 Cr	1305	9,195 / sq ft	Ready to move

```
df.dtypes
```

```
Unnamed: 0      int64
title           object
price           object
size            object
price_per_sqft  object
status          object
dtype: object
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74208 entries, 0 to 74207
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0             74208 non-null  int64
1   title                  74208 non-null  object
2   price                  74208 non-null  object
3   size                   73612 non-null  object
4   price_per_sqft         58445 non-null  object
5   status                 45372 non-null  object
dtypes: int64(1), object(5)
memory usage: 3.4+ MB
```

```
df.shape
```

```
(74208, 6)
```

```
df.drop_duplicates()
```

	Unnamed: 0	title	price	size	price_per_sqft	status	
0	0	3 BHK Apartment in CasaGrand Casagrand Meridian	1.18 Cr	2037	5,793 / sq ft	Under Construction	
1	1	2 BHK Apartment in Shree Nandana Elite	57.56 L	1240	4,642 / sq ft	Under Construction	
2	2	Residential Plot in Saroj Whispering Winds	32.3 L	2000	1,615 / sq ft	New	
3	3	3 BHK Apartment in Navami Landmaark	1.33 Cr	1641	8,133 / sq ft	Under Construction	
4	4	2 BHK Apartment in V Venture EVA	55.32 L	957	5,787 / sq ft	Ready to move	
...	
74203	74202	5 Bedroom Built-up Area in Whitefield	23100000	3453	6689.834926	NaN	
74204	74203	4 BHK Super built-up Area in Richards Town	40000000	3600	11111.111111	NaN	
74205	74204	2 BHK Built-up Area in Raja Rajeshwari Nagar	6000000	1141	5258.545136	NaN	
74206	74205	4 BHK Super built-up Area in Padmanabhanagar	48800000	4689	10407.336319	NaN	
74207	74206	1 BHK Super built-up Area in Doddathoguru	1700000	550	3090.909091	NaN	

74208 rows × 6 columns

▼ Forming a new column and adding the number bedroom into it from the 'title' column

```
df['space']=df['title'].str.split().str[:2].str.join(' ')
df.head()
```

	Unnamed: 0	title	price	size	price_per_sqft	status	space	
0	0	3 BHK Apartment in CasaGrand Casagrand Meridian	1.18 Cr	2037	5,793 / sq ft	Under Construction	3 BHK	
1	1	2 BHK Apartment in Shree Nandana Elite	57.56 L	1240	4,642 / sq ft	Under Construction	2 BHK	
2	2	Residential Plot in Saroj Whispering Winds	32.3 L	2000	1,615 / sq ft	New	Residential Plot	
3	3	3 BHK Apartment in Navami Landmaark	1.33 Cr	1641	8,133 / sq ft	Under Construction	3 BHK	
4	4	2 BHK Apartment in V Venture EVA	55.32 L	957	5,787 / sq ft	Ready to move	2 BHK	

▼ Drop unnecessary columns

```
df.drop(['Unnamed: 0', 'title', 'status'], inplace=True, axis=1)
```

▼ Removing irrelevant characters and units from the features

```
df['price_per_sqft']=df['price_per_sqft'].str.replace('/ sq ft','')
df['price_per_sqft']=df['price_per_sqft'].str.replace(',','') #5,751
df['price_per_sqft']=df['price_per_sqft'].str.replace('per sqft','')
df['price_per_sqft']=df['price_per_sqft'].str.replace('₹','')
df['size']=df['size'].str.replace('sqft','')
```

```
df['size']=df['size'].str.replace('sqm','')
df.head()
```

	price	size	price_per_sqft	space	
0	1.18 Cr	2037	5793	3 BHK	
1	57.56 L	1240	4642	2 BHK	
2	32.3 L	2000	1615	Residential Plot	
3	1.33 Cr	1641	8133	3 BHK	
4	55.32 L	957	5787	2 BHK	

```
df['space'].value_counts()
```

```
2 BHK      18655
3 BHK      18547
Residential Plot  11355
4 BHK       4055
1 BHK       1944
...
Ganapathipura Plot    1
Bilekahalli 6+        1
Seegehalli 3          1
Nelamangala 2         1
18 Bedroom            1
Name: space, Length: 2059, dtype: int64
```

▼ Replacing values in space column other than 'BHK' and 'Bedroom' into 'plot'

```
df['space']=df['space'].map(lambda x: x if 'BHK' in x or 'Bedroom' in x else 'plot' )
df.head()
```

```
# def take(sp):
#     if 'BHK' in sp:
#         pass
#     elif 'Bedroom' in sp:
#         pass
#     else:
#         sp='plot'
#     return sp

# df['space']=df['space'].apply(take)
```

	price	size	price_per_sqft	space	
0	1.18 Cr	2037	5793	3 BHK	
1	57.56 L	1240	4642	2 BHK	
2	32.3 L	2000	1615	plot	
3	1.33 Cr	1641	8133	3 BHK	
4	55.32 L	957	5787	2 BHK	

```
df['space'].value_counts()
```

```
plot      27101
2 BHK     18655
3 BHK     18547
4 BHK      4055
1 BHK      1944
4 Bedroom    813
5 BHK        662
3 Bedroom    547
2 Bedroom    329
5 Bedroom    297
6 BHK        246
6 Bedroom    191
7 BHK        154
10 BHK       145
1 Bedroom    105
8 BHK         99
8 Bedroom     84
7 Bedroom     83
9 BHK         60
9 Bedroom     46
10 Bedroom    12
11 BHK         7
12 BHK         7
```

```

15 BHK          6
13 BHK          4
11 Bedroom     2
27 BHK          1
19 BHK          1
16 BHK          1
43 Bedroom     1
14 BHK          1
12 Bedroom     1
18 Bedroom     1
Name: space, dtype: int64

```

Removing rows having plot value in space column

```

df.drop(df[df['space'] == 'plot'].index, inplace=True)
df.reset_index(drop=True,inplace=True)
df.head()

```

	price	size	price_per_sqft	space
0	1.18 Cr	2037	5793	3 BHK
1	57.56 L	1240	4642	2 BHK
2	1.33 Cr	1641	8133	3 BHK
3	55.32 L	957	5787	2 BHK
4	83.47 L	1575	5300	3 BHK

Removing 'BHK' and 'Bedroom' from the space column

```

df['space']=df['space'].str.replace('Bedroom','')
df['space']=df['space'].str.replace('BHK','')
df.head()

```

	price	size	price_per_sqft	space
0	1.18 Cr	2037	5793	3
1	57.56 L	1240	4642	2
2	1.33 Cr	1641	8133	3
3	55.32 L	957	5787	2
4	83.47 L	1575	5300	3

```

df.shape

```

(47107, 4)

```

df['space'].value_counts()

```

```

3      19094
2      18984
4       4868
1       2049
5        959
6        437
7         237
8         183
10        157
9         106
11          9
12          8
15          6
13          4
27          1
19          1
16          1
43          1
14          1
18          1
Name: space, dtype: int64

```

```

df.isna().sum()

```

```
price          0
size          13677
price_per_sqft 13205
space          0
dtype: int64
```

```
df.dropna(subset=['price_per_sqft'],inplace=True)
```

```
df.isna().sum()
```

```
price          0
size          546
price_per_sqft 0
space          0
dtype: int64
```

```
df.dtypes
```

```
price          object
size          object
price_per_sqft object
space          object
dtype: object
```

▼ string replace

```
df['price']=df['price'].str.replace('₹','')
```

▼ Converting price into numbers

```
def convert_prize(cash):
    if 'Cr' in cash:
        return float(cash.replace(' Cr',''))*10000000
    elif 'Lac' in cash:
        return float(cash.replace(' Lac',''))*100000
    elif 'L' in cash:
        return float(cash.replace(' L',''))*100000
    else:
        return float(cash)
```

```
df['price']=df['price'].apply(convert_prize)
```

```
df['price'].unique()
```

```
array([11800000.      , 5756000.      , 13300000.      , ...,
        64400000.00000001, 9710000.      , 9160000.      ],
      ])
```

▼ Datatype changing

```
df['size']=df['size'].astype(float)
df['price_per_sqft']=df['price_per_sqft'].astype(float)
df['space']=df['space'].astype(int)
df['price']=df['price'].astype(int)
```

```
df.isna().sum()
```


```
price          0
size          546
price_per_sqft 0
space          0
dtype: int64
```

```
df.shape
```

```
(33902, 4)
```

▼ Dealing missing values

```
sns.distplot(df['size'])
```



```
# distribution is not normal so median is used to fill the missing value  
df['size']=df['size'].fillna(df['size'].median())
```



```
df.isna().sum()
```

```
price      0  
size       0  
price_per_sqft  0  
space      0  
dtype: int64
```

```
df.dtypes
```

```
price      int64  
size      float64  
price_per_sqft  float64  
space      int64  
dtype: object
```


```
df.head()
```

	price	size	price_per_sqft	space	
0	11800000	2037.0	5793.0	3	
1	5756000	1240.0	4642.0	2	
2	13300000	1641.0	8133.0	3	
3	5532000	957.0	5787.0	2	
4	8347000	1575.0	5300.0	3	

▼ Graphical representation of the Data

```
sns.pairplot(df)
```

<seaborn.axisgrid.PairGrid at 0x7c3c2dec4610>




```
sns.heatmap(df.corr(),annot=True)
```

<Axes: >



```
sns.countplot(x='space',data=df)
```


<Axes: xlabel='space', ylabel='count'>




```
x_axis=['size', 'price_per_sqft', 'space']
y_axis=df['price']

for i in x_axis:
    print(sns.regplot(x=df[i],y=y_axis))
plt.show()
```



Axes(0.125,0.11;0.775x0.77)




Axes(0.125,0.11;0.775x0.77)



Axes(0.125,0.11;0.775x0.77)




sns.boxplot(df)




**there is outliers in price column **

```
sns.boxplot(df['price_per_sqft'])
```




```
sns.boxplot(df['size'])
```



```
sns.boxplot(df['space'])
```


<Axes: >



before removing outliers

```
sns.boxplot(df['price'])
```

<Axes: >



▼ dealing outliers with mean iqr - method

```
# interquartile range
ds=['size', 'price_per_sqft','price']
for i in ds:
    q1=df[i].quantile(0.25)
    q3=df[i].quantile(0.75)
    iqr=q3-q1
    lower=q1-(iqr*1.5)
    upper=q3+(iqr*1.5)

    df[i] = df[i].apply(lambda x: x if lower <= x <= upper else df[i].mean())
```

after dealing outliers

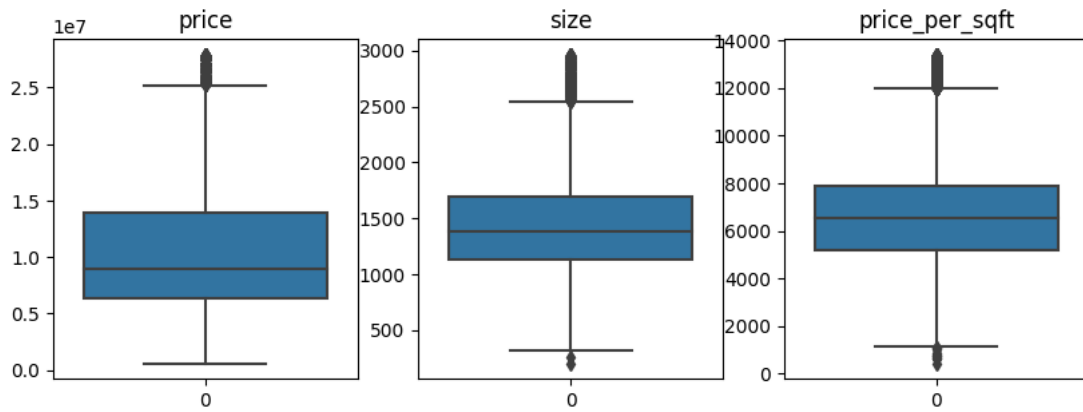
```
plt.figure(figsize=(10,3.3))
```

```
plt.subplot(1,3,1)
sns.boxplot(df['price'])
plt.title('price')
```

```
plt.subplot(1,3,2)
sns.boxplot(df['size'])
plt.title('size')
```

```
plt.subplot(1,3,3)
```

```
sns.boxplot(df['price_per_sqft'])
plt.title('price_per_sqft')
plt.show()
```



```
df.dtypes
```

```
price          float64
size           float64
price_per_sqft float64
space          int64
dtype: object
```

```
x=df.drop(['price'],axis=1).astype(int)
y=df['price'].astype(int)
```

▼ Feature selection using **chi_square test**

```
# from sklearn.feature_selection import SelectKBest,chi2

# chi=SelectKBest(chi2,k=3)
# best=chi.fit_transform(x,y)
#print(best.shape)
# x_inde=chi.get_support(indices=True)
# print(df.columns[x_inde])

# x_chi=df.drop(['space','price'],axis=1)
# x_chi.dtypes
```

```
from sklearn.model_selection import train_test_split
xtr,xts,ytr,yts=train_test_split(x,y,random_state=42,test_size=0.30)
```

```
from sklearn.preprocessing import StandardScaler
std=StandardScaler()
std.fit(xtr)
xtr=std.transform(xtr)
xts=std.transform(xts)
```

▼ **Model creation**

```
from sklearn.linear_model import LinearRegression
model=LinearRegression()
```


▼ Hyperparameter tuning for linear Regression

```
# to get the default values for the parameters
model.get_params()
```

```
{'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'positive': False}
```

```
from sklearn.model_selection import GridSearchCV

parameter={'copy_X': [True,False], 'fit_intercept': [True,False], 'n_jobs': [None,1,5,7,6], 'positive':[True, False]}
gsv=GridSearchCV(model,parameter,cv=10,scoring='accuracy')
gsv.fit(xtr,ytr)
```



```
gsv.best_params_

{'copy_X': True, 'fit_intercept': True, 'n_jobs': None, 'positive': True}
```

Multiple linear regression model creation

```
from sklearn.linear_model import LinearRegression
modell=LinearRegression(positive=True)
modell.fit(xtr,ytr)
ypr=modell.predict(xts)
```


```
df1=pd.DataFrame({'actual':yts,'predicted':ypr,'difference':yts-ypr})
df1
```

	actual	predicted	difference	
18397	13895895	1.369176e+07	2.041323e+05	
24636	13895895	2.244445e+07	-8.548554e+06	
5749	13895895	1.557182e+07	-1.675923e+06	
21115	13895895	2.003561e+07	-6.139714e+06	
1402	6400000	6.553946e+06	-1.539456e+05	
...	
19495	12500000	1.329248e+07	-7.924791e+05	
25818	16500000	1.560057e+07	8.994319e+05	
10202	2610000	1.686807e+06	9.231928e+05	
2984	11500000	1.274146e+07	-1.241461e+06	
8739	8769000	9.790120e+06	-1.021120e+06	

10171 rows × 3 columns

```
plt.figure(figsize=(8, 6))
plt.scatter(df1['actual'], df1['predicted'], marker='o', color='blue', label='Actual vs. Predicted')
plt.plot([min(df1['actual']), max(df1['actual'])], [min(df1['actual']), max(df1['actual'])], linestyle='--', color='red', label='Perfect')
plt.legend()
```

<matplotlib.legend.Legend at 0x7c3c276c7df0>



```
print('slope is ')
print(list(zip(x,model1.coef_)))
print('constant is ',model1.intercept_)
```

```
slope is
[('size', 2381157.806194569), ('price_per_sqft', 2703215.8259997703), ('space', 1207912.0173615823)]
constant is 10436601.962454174
```

```
from sklearn.metrics import r2_score,mean_absolute_percentage_error,mean_squared_error
```

```
r0=r2_score(yts,ypr)
print('r2 score',r2_score(yts,ypr))
print('maep ',mean_absolute_percentage_error(ypr,yts))
```

```
r2 score 0.7565905904339756
maep 0.16570523266452494
```

▼ Polynomial Regression

```
from sklearn.preprocessing import PolynomialFeatures
poly=PolynomialFeatures(degree=3)
poly.fit(x,y)
xply=poly.fit_transform(x)

xtrp,xtsp,ytrp,ytsp=train_test_split(xply,y,random_state=42,test_size=0.30)

model2=LinearRegression()
model2.fit(xtrp,ytrp)
yp=model2.predict(xtsp)
```

```
xply.shape
```

```
(33902, 20)
```

```
r1=r2_score(ytsp,yp)
print('r2 score',r1)
print('maep ',mean_absolute_percentage_error(yp,ytsp))
```

```
r2 score 0.8134061502988745
maep 0.11769392929248679
```

▼ Decision tree algorithm

```
from sklearn.tree import DecisionTreeRegressor
dec=DecisionTreeRegressor()
dec.fit(xtr,ytr)
ypr1=dec.predict(xts)

r2=r2_score(yts,ypr1)
print('r2 score',r2)
print('maep ',mean_absolute_percentage_error(ypr1,yts))
```

```
r2 score 0.9223807263497109
maep 0.03143447676581806
```

▼ Random forest algorithm

```
from sklearn.ensemble import RandomForestRegressor
random=RandomForestRegressor()
random.fit(xtr,ytr)
ypr2=random.predict(xts)

r3=r2_score(yts,ypr2)
print('r2 score',r3)
print('maep ',mean_absolute_percentage_error(ypr2,yts))
```

```
r2 score 0.9412874225012221
maep 0.03316465026566907
```

▼ Ridge Regression

```
from sklearn.linear_model import Ridge
rdg=Ridge(alpha=2)
rdg.fit(xtr,ytr)
yr=rdg.predict(xts)

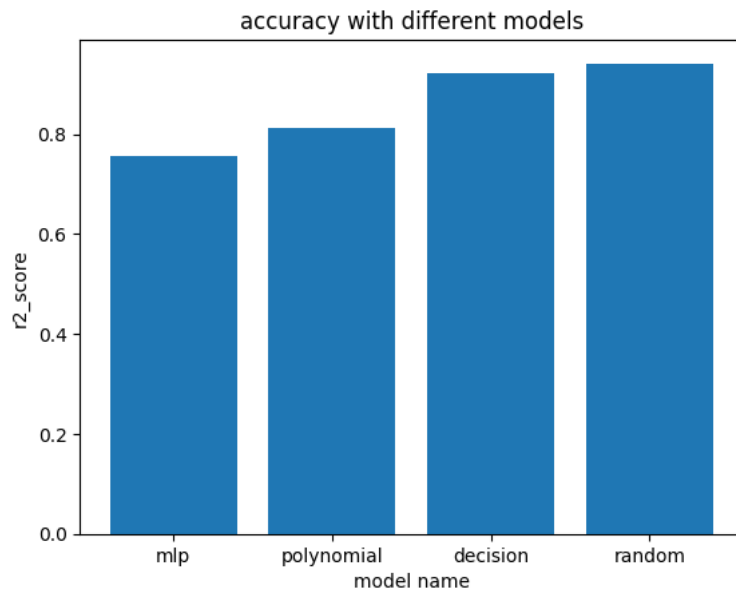
r4=r2_score(yts,yr)
print('r2 score',r4)
print('maep ',mean_absolute_percentage_error(yr,yts))
```

```
r2 score 0.75659235954779
maep 0.16659618732812337
```

```
al=['mlp','polynomial','decision','random']
result=[r0,r1,r2,r3]
```

```
plt.bar(al,result)
plt.xlabel('model name')
plt.ylabel('r2_score')
plt.title('accuracy with different models')
```

```
Text(0.5, 1.0, 'accuracy with different models')
```



▼ Pickling

```
import pickle

with open('random_forest.pickle', 'wb') as dump_var:
    pickle.dump(random, dump_var)

pickle_in = open('random_forest.pickle', 'rb')
pickle_clf = pickle.load(pickle_in)

accuracy_pkl = pickle_clf.score(xts,yts)
accuracy_pkl
```

0.9412874225012221